

The AppleCore Language Specification, v1.0

Robert L. Bocchino Jr.
Pittsburgh, PA

March 23, 2013

Contents

1	Introduction and Rationale	2
2	Lexical Structure	3
2.1	White Space	3
2.2	Identifiers	3
2.3	Keywords	4
2.4	Constants	4
2.5	Symbols	4
2.6	Comments	5
3	Syntax	5
3.1	Source Files	6
3.2	Declarations	6
3.3	Statements	7
3.4	Expressions	8
3.5	Types	9
4	Semantic Checking	9
4.1	Representation of Values	9
4.2	Imported Declarations	10
4.3	Attribution	10
4.4	Constant-Value Expressions	11
4.5	Size, Signedness, and Addresses	11
4.6	Number of Function Arguments	13
4.7	Semantic LValues	14

5	Code Generation	14
5.1	Source Files	14
5.2	Global Declarations	14
5.3	Program Stack	15
5.4	Function Bodies	16
5.5	Statements	16
5.6	Expressions	17
5.7	Address Computation for Semantic LValues	18
5.8	Size Adjustment	19

1 Introduction and Rationale

AppleCore is a C-like programming language for the Apple II series of computers. A cross-compiler from Mac OS X to the Apple II is currently available.

The idea of AppleCore is to provide a “core language for programming the Apple II.” The motivation is as follows:

1. Programming the Apple II is both nostalgic and just plain fun.
2. The Apple II is a small machine. Therefore it needs a lightweight programming model (i.e., without complicated runtime overhead in space or time). The programming model must also provide (1) tight control over the layout of generated code and (2) easy integration with assembly-language support routines such as those provided by the system monitor or Applesoft.
3. The native programming models on the Apple II are inadequate. Assembly language is powerful but painfully low-level. Applesoft, on the other hand, is a disaster. It lacks even the most basic abstraction required of a sane programming environment, including scoped loops, scoped functions with local variables, and variable names longer than *two letters*! Applesoft does have nice support for numeric and string input and output, so if you just want to write short programs that read and write things, it’s handy. But for any but the tiniest programs, Applesoft code becomes an unreadable mess. In terms of control abstraction Applesoft is much worse than assembly language, because it forces you to branch to numbered lines, whereas in assembly language you can JMP and JSR to symbolic labels with meaningful names.

The goal of AppleCore is to rectify this situation by providing a “low-level high-level” language (at about the same level of abstraction as C, but without structs) that’s both useful and fun to program with. AppleCore is quite expressive for low-level programming: it supports all the usual assembly-language techniques such as symbolic constants, calling addresses as subroutines, passing and storing subroutines as addresses, laying out data tables and jump tables, etc. However, AppleCore also provides higher-level abstractions such as loops, function prototypes, multi-byte variables and arithmetic, and multi-byte loads and stores to memory, making it much easier to program than assembly. The dollar notation for hexadecimal numbers is supported, so addresses, masks, ASCII codes, etc. can be expressed in a convenient, Apple II-friendly style.

AppleCore supports separate compilation, so libraries can be statically linked in as needed; together with function abstraction, this capability provides a simple but powerful kind of extensibility. For example, whereas Applesoft specifies many built-in keywords pertaining to graphics operations, AppleCore has a very small set of built-in keywords. Capabilities pertaining to graphics, IO, sound, etc., are provided as libraries.

The language also has several unusual features that are perhaps interesting as a matter of language design. First, AppleCore natively supports integer arithmetic (both signed and unsigned) on fixed-size variables of up to 255 bytes. For example, you can declare and initialize variables A and B, each of size 16 bytes unsigned, and then the expression A+B also has size 16, and it represents the 16-byte unsigned addition of A and B in the obvious way. (Variable-sized arithmetic, i.e., where the size is not known at compile time, is supported by library functions.) Second, a simple type system keeps track of size and signedness. The programmer specifies the size (as a compile-time constant-value expression) and signedness of each variable and function return value in the program, and the compiler infers the rest. Third, *register expressions* allow AppleCore programs to read and write 6502 registers. Register access is essential for low-level routines, e.g., a character printing routine that loads a value into the accumulator and then calls MON_COUT, and register expressions provide a much cleaner solution than inline assembly. Inline assembly is not supported, but it's easy to link against handwritten assembly and call it from within AppleCore code.

The main things that AppleCore lacks are heap memory management (stack allocation is supported) and floating-point computation. I've left out both of these features to keep things very simple, and because they aren't needed for the programming projects I am interested in. Either or both could be added without difficulty. One can also carry out FP computations by encoding FP literals as string data and calling into the Applesoft FP routines to convert from strings to FP numbers and back, and to compute using the FP numbers.

2 Lexical Structure

Before parsing a source file, the source text is separated into *tokens*. There are four classes of tokens: identifiers, keywords, constants, and symbols.

2.1 White Space

White space consists of any sequence of the following characters: space (ASCII SP, value \$20), newline (ASCII NL, value \$0A), carriage return (ASCII CR, value \$0D), and horizontal tab (ASCII HT, value \$09). Whitespace is syntactically relevant only in the following ways:

1. *Separation of tokens.* When a whitespace character or end of file appears immediately after a non-whitespace character, that signifies the end of a token.
2. *End of line.* An end-of-line sequence (EOL) is a comment terminator (see Section 2.6). As usual, the definition of EOL is platform-dependent: on the Apple II it is CR, on UNIX it is NL, and on Windows it is CR followed by NL.

The lexer may also count source lines (using EOL) to provide line numbers for error messages.

2.2 Identifiers

An identifier is a sequence of letters, underscore character, and digits that is not a keyword. The first character must be a letter. Upper- and lowercase letters are distinct. A single identifier may contain at most 32 characters.

2.3 Keywords

The following sequences of characters are reserved for use as keywords and may not be used as identifiers:

AND	CONST	DATA	DECR
ELSE	FN	IF	INCLUDE
INCR	NOT	OR	RETURN
SET	VAR	XOR	WHILE

Following Apple II tradition (mostly because the original Apple II had no support for lower-case letters), AppleCore keywords are uppercase. That means that the same words in lowercase (or a combination of upper- and lowercase) are not recognized as keywords, so those sequences of characters are available for use as identifiers.

2.4 Constants

There are three types of constants: integer constants, string constants, and character constants.

Integer constants. Integer constants may be written in decimal or hexadecimal form. A decimal integer constant consists of one or more decimal digits 0 through 9. For example, 1 and 123 are valid integer constants in decimal form. A hexadecimal integer constant consists of a dollar sign \$ followed by one or more hexadecimal digits 0 through 9 or A through F. For example, \$1, \$A, and \$00FF are valid integer constants in hexadecimal form.

Integer constants represent values with up to 255 bytes of precision (i.e., in the range 0 through $2^{255 \cdot 8} - 1$, inclusive, when interpreted as unsigned integers). It is a compile-time error to write an integer constant with a value larger than that.

String constants. A string constant is a sequence of characters enclosed in double quotes ("..."). A string constant represents a sequence of ASCII characters, one for each character appearing in the double-quotes, except that the character sequence \\$ has the special meaning described below.

Since not all ASCII characters have printable representations, an *escape sequence* may be used to represent an arbitrary ASCII value (printable or non-printable). An escape sequence consists of a backslash \ followed by a dollar sign \$ and two hexadecimal digits. The whole sequence represents the single character with the ASCII value given by the digits. For example, the string constant

"Hello, world!\\$0D"

represents a string consisting of the characters Hello, world! followed by a CR character. Similarly, the quote character " can be embedded in a string constant with the sequence \\$22.

Character constants. A character constant consists of a single-quote character ', followed by a printable ASCII character, followed by another single-quote character. It represents the ASCII value associated with the character. For example, the constant 'A' represents the value \$41.

2.5 Symbols

AppleCore uses the symbols shown in Figure 1, each of which is a separate token.

Symbol	Meaning
@	Denotes the address of a variable, or an address type
~	Denotes a 6502 register expression
*	Multiplication
/	Division
+	Addition
-	Negation (as unary operator); subtraction (as binary operator)
<<	Left shift
>>	Right shift
>=	Greater than or equal to
<=	Less than or equal to
>	Greater than
<	Less than
=	Equal to; assignment
(and)	Encloses function parameters and arguments, parenthesized expressions
{ and }	Encloses statement blocks
[and]	Address dereference
;	Terminates declarations and statements
:	Separates variable declaration from size
,	Separates function parameters and arguments; separates index from size in dereference
\	Signifies unterminated string data

Figure 1: Symbols used in AppleCore syntax.

2.6 Comments

The character # indicates a comment; all text to the next end of line (or end of file, if there is no end of line) are ignored by the lexer. Multi-line comments are indicated by preceding each line with #.

3 Syntax

The syntax description below uses the following conventions:

- The symbol * denotes zero or more instances of the entity preceding it.
- Italicized parentheses () group the enclosed symbols and do not denote program text.
- Italicized brackets [] signify that the enclosed symbols are optional (i.e., they may occur zero or one time). They do not denote program text.
- The nonterminal *identifier* stands for any identifier as defined in Section 2.2.
- The nonterminals *integer-const*, *string-const*, and *char-const* stand for integer, string, and character constants as defined in Section 2.4.
- Text and symbols in typewriter font (including non-italicized parentheses and brackets) denote literal program text.

3.1 Source Files

The basic syntactic unit of an AppleCore program is a *source file*, i.e., an input file presented to the AppleCore compiler for compilation. The compiler translates the source file into an assembly file which is then linked with other assembly files as discussed in Section 5.1 to form a complete executable program.

An AppleCore source file is given by zero or more declarations (Section 3.2):

$$\text{source-file} ::= \text{decl}^*$$

3.2 Declarations

A declaration is a constant declaration, a data declaration, a variable declaration, a function declaration, or an include declaration:

$$\text{decl} ::= \text{const-decl} \mid \text{data-decl} \mid \text{var-decl} \mid \text{fn-decl} \mid \text{include-decl}$$

Constant declarations. A constant declaration consists of the keyword CONST, an identifier, an optional expression, and a terminating semicolon:

$$\text{const-decl} ::= \text{CONST identifier} [\text{expr}] ;$$

Data declarations. A data declaration consists of the keyword DATA, an optional identifier representing a label for the data, an expression or a string constant, and a terminating semicolon. A backslash may optionally follow the string constant, indicating that the string is unterminated (see Section 5.2).

$$\text{data-decl} ::= \text{DATA} [\text{identifier}] (\text{expr} \mid (\text{string-const} [\backslash])) ;$$

Variable declarations. A variable declaration consists of the keyword VAR, an identifier representing the variable name, a type, an optional initializer expression, and a terminating semicolon:

$$\text{var-decl} ::= \text{VAR identifier} : \text{type} [= \text{expr}] ;$$

Function declarations. A function declaration consists of the keyword FN, an identifier representing the function name, the function parameters enclosed in parentheses, an optional return type, and the function body:

$$\text{fn-decl} ::= \text{FN identifier} (\text{fn-params}) [\text{: type}] \text{fn-body}$$

The function parameters are a comma-separated list of zero or more parameters:

$$\text{fn-params} ::= [\text{fn-param} (, \text{fn-param})^*]$$

A parameter consists of an identifier representing the parameter name and type:

$$\text{fn-param} ::= \text{identifier} : \text{type}$$

A function body is either zero or more variable declarations and statements enclosed in braces, or a semicolon indicating an externally defined function:

$$fn-body ::= \{ var-decl^* stmt^* \} \mid ;$$

Include declarations. An include declaration consists of the keyword INCLUDE, a string constant, and a terminating semicolon:

$$include-decl ::= INCLUDE string-const ;$$

3.3 Statements

A statement is an if statement, a while statement, a set statement, a call statement, and increment statement, a decrement statement, return statement, or a block statement:

$$stmt ::= if-stmt \mid while-stmt \mid set-stmt \mid call-stmt \mid incr-stmt \mid decr-stmt \mid return-stmt \mid block-stmt$$

If statements. An if statement consists of the keyword IF, a conditional expression enclosed in parentheses, a statement to execute if the condition is true, and optionally the keyword ELSE followed by a statement to execute if the condition is false:

$$if-stmt ::= IF (expr) stmt [ELSE stmt]$$

While statements. A while statement consists of the keyword WHILE, a test expression enclosed in parentheses, and a statement to execute as long as the condition is true:

$$while-stmt ::= WHILE (expr) stmt$$

Set statements. A set statement consists of the keyword SET, an expression, an equals sign, a right-hand-side expression, and a terminating semicolon:

$$set-stmt ::= SET expr = expr ;$$

The left-hand side expression may not be a binary operation =.

Call statements. A call statement consists of a call expression followed by a terminating semicolon:

$$call-stmt ::= call-expr ;$$

Increment statements. An increment statement consists of the keyword INCR, an expression, and a terminating semicolon:

$$incr-stmt ::= INCR expr ;$$

Decrement statements. A decrement statement consists of the keyword DECR, an expression, and a terminating semicolon:

$$\text{decr-stmt} ::= \text{DECR } \text{expr} ;$$

Return statements. A return statement consists of the keyword RETURN, an optional expression, and a terminating semicolon:

$$\text{return-stmt} ::= \text{RETURN } [\text{expr}] ;$$

Block statements. A block statement consists of zero or more statements enclosed in braces:

$$\text{block-stmt} ::= \{ \text{stmt}^* \}$$

3.4 Expressions

An expression is an identifier, an indexed expression, a register expression, a numeric constant, a call expression, a binary operation expression, a unary operation expression, a parentheses expression, or a typed expression.

$$\begin{aligned} \text{expr} ::= & \text{identifier} \mid \text{indexed-expr} \mid \text{register-expr} \mid \text{numeric-const} \mid \text{call-expr} \mid \\ & \text{binop-expr} \mid \text{unop-expr} \mid \text{parens-expr} \mid \text{typed-expr} \end{aligned}$$

Except for binary expressions as specified below, any ambiguity in parsing is resolved using right recursion. For example, @F() is parsed as @(F()) and not (@F)().

Numeric constants. A numeric constant is an integer or character constant:

$$\text{numeric-const} ::= \text{integer-const} \mid \text{char-const}$$

Call expressions. A call expression consists of an expression followed by an argument list enclosed in parentheses:

$$\text{call-expr} ::= \text{expr} ([\text{expr} (, \text{expr})^*])$$

Indexed expressions. An indexed expression consists of a base expression, an offset expression, and a type:

$$\text{indexed-expr} ::= \text{expr} [\text{expr} , \text{type}]$$

Register expressions. A register expression consists of a caret character followed by a 6502 register name:

$$\text{register-expr} ::= \wedge (\text{A} \mid \text{X} \mid \text{Y} \mid \text{P} \mid \text{S})$$

Binary operation expressions. A binary operation expression consists of a left-hand-side expression, a binary operator, and a right-hand-side expression:

$$\begin{aligned} \text{binop-expr} &::= \text{expr binop expr} \\ \text{binop} &::= > \mid < \mid <= \mid >= \mid \text{AND} \mid \text{OR} \mid \text{XOR} \mid + \mid - \mid * \mid / \mid << \mid >> \mid = \end{aligned}$$

Parsing of binary operation expressions is disambiguated using the following precedence rules:

1. In any sequence *expr-1 binop-1 expr-2 binop-2 expr-3*, the implied parentheses go around the first operation unless the second operation has a strictly higher precedence. There are five levels of precedence, from highest to lowest: (1) << and >>; (2) * and \; (3) + and -; (4) =, >, <, <=, and >=; and (5) AND, OR, and XOR.
2. Binary operators bind more weakly than any other expression combinators. For example, *A+B()* is parsed as *A+(B())* and not *(A+B)()*; and *-5+3* is parsed as *(-5)+3* and not *-(5+3)*.

Unary operation expressions. A unary operation expression consists of a unary operator followed by an expression:

$$\textit{unop-expr} ::= \textit{unop expr}$$

$$\textit{unop} ::= @ \mid \text{NOT} \mid -$$

Parentheses expressions. A parentheses expression is an expression surrounded by parentheses:

$$\textit{parens-expr} ::= (\textit{expr})$$

Typed expressions. A typed expression is an expression with an explicit type:

$$\textit{typed-expr} ::= \textit{expr} : \textit{type}$$

3.5 Types

A type is an unsigned size, a signed size, or an address type:

$$\textit{type} ::= \textit{expr} [\textit{S}] \mid @$$

The size expression may not be a binary operation =.

4 Semantic Checking

The compiler checks every source file for conformance to the semantic rules stated below. A semantically invalid program (i.e., one that violates one of these rules) generates a compile-time error and causes compilation to halt. Only semantically valid files go on to code generation as described in Section 5.

4.1 Representation of Values

The AppleCore language supports two types of values: integer values and Boolean values.

Integer values. Integer values are represented in AppleCore in the ordinary way for the Apple II:

1. Unsigned single-byte values are represented by treating the 8 bits, low to high, as the 8 digits of a binary number.

2. Multi-byte values are stored in little-endian order.
3. Signed values are stored in two's complement representation (i.e., by constructing the absolute value of the number, flipping the bits, and adding one).

Boolean values. The AppleCore language has no special Boolean type; as in C, ordinary integers serve as Boolean values. However, AppleCore differs from C in what values it treats as “true” and “false.” In AppleCore, only the lowest-order bit is relevant in testing the Boolean value of an integer: integers whose low bit is 1 (i.e., odd integers) all count as “true,” while integers whose low bit is 0 (i.e., even integers) all count as “false.” This mechanism allows bitwise NOT to function as logical NOT, without any need for a separate logical negation operator, and that keeps the language simple. Also, the constants 1 and 0 can function as “true” and “false.” These names and values are not built into the language, but one can easily declare constants TRUE and FALSE with the values 1 and 0.

4.2 Imported Declarations

Before performing any semantic checks, the compiler imports zero or more declarations specified by the user, usually as a compiler option. The option must allow the user to specify (1) the files from which to read the declarations and (2) the order in which to read the files.

Except for include declarations, each declaration so encountered is logically prepended to the source file under translation, in the order encountered, subject to the following rules:

1. Constant declarations are prepended as if they were literally included in the source file.
2. Function declarations are prepended as if they were literally included in the source file, but with a semicolon instead of the actual body (if any).
3. Variable and data declarations are prepended storing their name and size only; no variable initializer or data information is included for attribution (Section 4.3), and these declarations generate no code.

External include declarations are not imported, because they do not contain any information that the compiler needs for attribution or code generation of the source file.

4.3 Attribution

The first semantic check performed by the compiler is to match each identifier appearing in the source file with its corresponding definition; this step is called *attribution*. A *definition* is a constant declaration, data declaration, variable declaration, function parameter declaration, or function declaration. A *use* is any identifier appearing in the source file not as the identifier of a definition. Identifiers appearing in imported variable and data declarations (Section 4.2) do not appear in the source file for purposes of this rule.

The compiler matches definitions with uses in the following way:

1. At global scope (i.e., outside of any function body), for each use find the corresponding definition with the same identifier at global scope. If there is no corresponding definition, or if any two definitions use the same identifier, then report an error.
2. At function scope (i.e., inside a function body), for each use find the corresponding definition with the same identifier in the function scope, which consists of the global scope together with all function parameters and local variables declared in the function. If there is no corresponding definition, or if any two definitions in the function scope use the same identifier, then report an error.

Constant declarations must precede their uses in the source file. Otherwise, declarations need not precede their uses.

4.4 Constant-Value Expressions

The term “constant-value expression” refers to either a compile-time constant-value expression or an assembly-time constant-value expression.

Compile-time constant-value expressions. A compile-time constant-value expression is one of the following:

1. A numeric constant.
2. An identifier corresponding to a constant declaration.
3. A binary operation expression, where each operand is a compile-time constant-value expression.
4. A negation (−) or NOT unary operation expression, where the operand is a compile-time constant-value expression.
5. A parentheses expression, where the expression inside the parentheses is a compile-time constant-value expression.
6. A typed expression, where the value subexpression is a compile-time constant-value expression.

The following must be compile-time constant-value expressions: (1) Any expression appearing in a constant declaration (Section 3.2); and (2) any expression appearing in the size expression of a type (Section 3.5).

Assembly-time constant-value expressions. An assembly-time constant-value expression is one of the following:

1. A compile-time constant-value expression.
2. An identifier corresponding to a data declaration (useful for data tables).
3. An identifier corresponding to a function declaration (useful for jump tables).
4. An address expression `@identifier`, where *identifier* is an identifier corresponding to a global variable declaration (useful for creating a separate global symbol for the address of a global variable).

Any expression appearing in a data declaration or as the initializer of a global variable declaration (Section 3.2) must be an assembly-time constant-value expression.

4.5 Size, Signedness, and Addresses

Every expression appearing in the source file is given a *size* and a *signedness*. An expression may (or may not) *represent an address*. The address representation rules are designed to exclude a class of easy-to-make yet potentially disastrous errors such as dereferencing a variable’s contents instead of its address (for example, writing `X[0, 1]` instead of `(@X)[0, 1]`).

Size and signedness associated with type. The compiler assigns an size and signedness to each type appearing in the program as follows:

1. If the type has the form *expr*, then use the rules in Section 5.6 to evaluate *expr* to a numeric constant *size*, which must be between 1 and 255 inclusive. The type has size *size* and is unsigned.
2. If the type has the form *expr* S, then proceed as in case 1, except that the resulting type is signed.

3. The type @ has size 2 and is unsigned.

Computing the size and signedness of an expression. For a constant-value typed expression (Section 4.4), the size and signedness come from the type. For any other compile-time constant-value expression, the size and signedness come from the value produced by evaluating the expression as specified in Section 5.6, assuming that the operands to any binary expression have the maximum allowed size of 255. The size is the minimum number of bytes required to represent the value. The expression is signed if the value is less than zero, otherwise unsigned.

For any other expression, if the expression appears in a data declaration or global variable declaration, then size is 2 and the signedness is unsigned. Otherwise the size and signedness are computed as follows:

- *Identifiers.* The size and signedness of an identifier come from its definition (Section 4.3). (1) If the definition is a variable declaration or function parameter declaration, then the size and signedness come from the type. (2) If the definition is a data declaration or function declaration, then the size is 2, and the signedness is unsigned.
- *Indexed expressions.* The size and signedness of an indexed expression come from the *type* component of the expression.
- *Register expressions.* The size of a register expression is 1. It is unsigned.
- *Call expressions.* If the called expression (i.e., the expression before the first parenthesis) is anything other than an identifier corresponding to a function declaration, then the size is 0 and the signedness is unsigned. Otherwise the size and signedness come from the function declaration corresponding to the identifier: if the function has a return type, then the size and signedness come from the type; otherwise the size is 0 and the signedness is undefined.
- *Binary operation expressions.* For comparison operations, the size is 1 byte unsigned. For shift operations, the size and signedness come from the left operand. For all other operations, the size is the maximum of the sizes of its operands. If either of the operands is signed, then the result is signed. Otherwise, the result is unsigned.
- *Unary operation expressions.* The size and signedness of a unary operation expression come from the operand, except that (1) an address operation @ is two bytes unsigned; and (2) a negation operation is signed.
- *Parentheses expressions.* The size and signedness of a parentheses expression come from the enclosed expression.
- *Typed expressions.* The size and signedness of a typed expression come from its type.

Address expressions. The following expressions each represent an address:

- *Identifiers.* An identifier represents an address if its associated definition (1) is a constant declaration of size 1 or 2 unsigned; or (2) is a data declaration; or (3) is a function declaration; or (4) is a variable or function parameter declaration whose type is @.
- *Indexed expressions.* An indexed expression represents an address if its *type* component is @.
- *Call expressions.* A call expression represents an address if the called expression is an identifier whose associated definition is a function with return type @.
- *Binary expressions.* A shift or divide operation represents an address if its left operand does. A binary operation that is not a shift, divide, or compare represents an address if exactly one of its operands does.

- *Unary operation expressions.* An address operation @ represents an address. Otherwise, a unary operation represents an address if its operand does.
- *Parentheses expressions.* A parentheses expression represents an address if its subexpression does.
- *Typed expressions.* A typed expression represents an address if the type is @.

No other expression represents an address.

Size, signedness, and address requirements for expressions. The following requirements apply to all expressions, whether or not they are constant-value expressions:

- *Indexed expressions.* The base expression (before the first bracket) of an indexed expression must represent an address. The index expression (immediately after the first bracket) must have size 1 or 2 and be unsigned.
- *Call expressions.* (1) The called expression must represent an address. (2) Each argument expression must have nonzero size. (3) Each argument must satisfy the requirements for a valid address assignment (see below), treating the argument expression as the right-hand side of the assignment, and the corresponding function parameter as the left-hand side.
- *Binary operation expressions.* For shift operations, the right-hand expression must be 1 byte unsigned.

Variable declarations. If an initializer expression is present, then (1) the expression must have non-zero size; and (2) the expression must satisfy the requirements for a valid address assignment, treating the variable declaration itself as the left-hand side, and the initializer expression as the right-hand side.

Set statements. (1) The right-hand expression of a set statement must have nonzero size. (2) The assignment represented by the statement must satisfy the requirements for a valid address assignment.

Return statements. (1) No return statement may contain an expression of zero size. (2) If a function has nonzero return size, and the function has a body, then (a) the body must end with a return statement containing an expression, and (b) every return statement appearing in the body must contain an expression. (3) If a function has return size 0, then no return statement in the function body may contain an expression. (4) A return statement that contains an expression must satisfy the requirements for a valid address assignment, treating the contained expression as the right-hand side, and treating the left-hand side as a fresh variable whose type is the return type of the enclosing function.

Valid address assignments. (1) If the left-hand side of an assignment represents an address, then the right-hand side must also represent an address. (2) If the right-hand side of an assignment represents an address and is not a constant-value expression, then the left-hand side must also represent an address.

Function frame sizes. The *frame size* of a function declaration that contains a body equals the sum of the sizes of all the function parameters and local variables. No function may have a frame size of greater than 255 bytes.

4.6 Number of Function Arguments

At every call expression where the called expression is an identifier corresponding to a function declaration, the number of arguments must match the number of parameters given in the declaration.

4.7 Semantic LValues

Definition of semantic lvalue. A *semantic lvalue* is an expression that (1) is an identifier, register expression, or indexed expression; and (2) if it is an identifier, then it corresponds to a variable or function parameter declaration.

Requirement of semantic lvalues. (1) A semantic lvalue is required in the following places: (a) on the left-hand side of a set statement; (b) as the expression of an increment or decrement statement; and (c) as the operand of an address (@) unary operator. If an expression that is not a semantic lvalue appears in any of these places, then a compile-time error results. (2) A compile-time error results if a register expression appears as the operand of an address (@) operator.

5 Code Generation

5.1 Source Files

The AppleCore language is designed to support separate compilation. To do that, the compiler translates source files to assembly files in one of two modes: *top-level mode* and *include mode*. The mode must be specified at the time the source file is translated to assembly, usually as a compiler option.

Top-level mode. In top-level mode, the compiler translates the source file as follows:

1. Issue some code (or an assembler directive to include the code) that sets up the program stack (Section 5.3).
2. Issue code to transfer control to the first function with a body that appears in the file. There must be at least one such function, and it is a compile-time error if not.
3. Translate all global declarations in the source file as described in Section 5.2, in the order in which they appear in the source file. Include declarations may appear anywhere in the file, and they are translated to directives to include the corresponding assembly files.
4. Issue assembler directives to include the AppleCore support code (for example, the code needed to do arithmetic operations).

Include mode. In include mode, the compiler just translates the declarations in the source file as described in Section 5.2.

Final assembly and symbol resolution. A user program consists of exactly one file translated in top-level mode and zero or more files translated in include mode. The final step in compilation is to ask the assembler to assemble the file translated in top-level mode, which includes all the other files for assembly. This assembly also causes name resolution for externally defined names. It is the user's responsibility to ensure that externally-defined names are both available and uniquely defined; the compiler cannot check this. If a symbol used in one of the assembled files is not defined in any of the files (or if any symbol is multiply defined), then an assembler error will result.

5.2 Global Declarations

Global declarations are translated as stated below, except for imported declarations (Section 4.2), which generate no code. To the extent that translation produces actual bytes of code, data, or reserved storage, the bytes must be laid out in the order in which the translated constructs are encountered in the source file.

Constant declarations. If the constant declaration *const-decl* contains an expression *expr*, then the compiler evaluates *expr* to a numeric constant according to the rules in Section 5.6. In the rest of the source file where *const-decl* appears, the compiler replaces each instance of the identifier given in *const-decl* expression with the value so computed. The compiler does *not* generate a label or other assembly artifact associated with *const-decl*. In particular, the same identifier may appear in different constant declarations in different files that are part of the same final assembly (Section 5.1).

If the constant declaration contains no expression, then the name in the declaration must refer to the two-byte unsigned value associated with an assembler label during final assembly (Section 5.1).

Data declarations. (1) If an expression appears before the semicolon, then the compiler evaluates the expression to an assembler expression according to the rules in Section 5.6 and generates code to store the bytes of the constant in little-endian order. (2) If a string constant appears before the semicolon, then the compiler generates assembly code to store the bytes of the string constant in left-to-right order. (a) If no backslash \ follows the string constant, then the compiler adds a terminating NUL (\$00) character after the last byte of string data. (b) If a backslash follows the string constant, then no NUL character is added. (3) In any case, if an identifier appears after the keyword DATA, then the compiler associates the name with the lowest address in which data is stored.

Global variable declarations. The compiler sets aside a number of bytes of storage equal to the declared size of the variable and associates the identifier appearing after VAR with the lowest address of this storage. If an expression appears before the semicolon, then the compiler evaluates the expression to an assembler expression according to the rules stated in Section 5.6 and issues assembly code to fill in the storage set aside for the variable with the bytes of the numeric constant, in little-endian order, using size adjustment (if necessary) as stated in Section 5.8.

Function declarations. (1) A function declaration with no body does not correspond to any generated code; its signature is used only for checking call expressions that invoke the function. (2) A function declaration with a body causes code to be generated as stated in Section 5.4.

Include declarations. The compiler translates an include declaration into an assembler directive to include the file named in the string constant.

5.3 Program Stack

This specification requires that the implementation provide three logical stacks for use at runtime, collectively referred to as the *program stack*:

1. A C-style *call stack*, such that a new frame appears at the top of the stack for each function call, and the frame is popped at the end of the call. This stack is used in generating function call code, as specified in Section 5.4.
2. An *expression stack* for use during expression evaluation as specified in Section 5.6.
3. An *allocation stack* that supports an ALLOCATE library function that works as follows. (1) ALLOCATE is only ever called when no temporary expression result is on the expression stack; otherwise the results are undefined. (2) When ALLOCATE is called, it is passed a two-byte size value *s*. (3) (a) If there is enough space on the stack, then ALLOCATE returns a pointer to *s* bytes on the stack. (b) If there are not *s* bytes remaining on the stack, then the results are undefined. (4) The *s* bytes pointed to by the result of the ALLOCATE call remain valid (i.e., are not modified except by user-written code) during the lifetime of the call stack frame for the function in which the ALLOCATE call occurred. (5) At the end of the function in which the ALLOCATE call occurred, the memory is deallocated and may be reused as program stack storage.

The compiler may provide this stack functionality in any way that is feasible. Typically the same physical stack would be used for all three logical stacks (and the logical stacks have been specified so that this is possible), but this is not a requirement. Also, typically the generated code would maintain its own two-byte stack pointer, because the native Apple II stack located at addresses \$200 through \$2FF is too small to serve as the program stack for this language. However, this specification does not mandate any particular stack size; the actual available stack size depends on both the compiler implementation and the program memory requirements (for example, whether the graphics and/or language card areas are available for program heap and stack data).

5.4 Function Bodies

The AppleCore language definition assumes a C-style function call stack as specified in Section 5.3. For each function declaration that has a body, the compiler does the following:

- Generate code to store the callee's frame pointer on the call stack.
- Generate code to reserve slots on the call stack for the function parameters, local variables, and saved registers. The saved registers are all the registers appearing in register expressions in the function body.
- For each local variable declaration that includes an initializer expression, generate code to evaluate the expression according to the rules in Section 5.6 and assign the resulting value into the stack slot for the variable, using size adjustment (if necessary) as stated in Section 5.8.
- For each statement of the function body, generate code as stated in Section 5.5. The statements of a function body are executed in sequence, in the order in which they appear in the program text.
- If the last statement executed in the function is not a return statement, then generate code to pop the frame off the stack and restore the parent frame.

This calling convention is designed so that from the point of view of the caller, a call to an AppleCore function that takes no arguments and returns no result is indistinguishable from a JSR to a non-AppleCore function (for example, in the system monitor). That way, non-AppleCore assembly language functions can be called in the same way as AppleCore functions. In particular, if *V* is a variable, then a call expression *V*() has the same meaning regardless of what kind of function address is stored in *V* (which is not generally known at compile time). Also, a source file with a function declaration `FN F00()`; can be assembled with any assembly file containing a function with label `F00`, regardless of whether the code at label `F00` uses the AppleCore conventions for managing the stack.

5.5 Statements

The compiler generates code to do the following for each program statement.

If statements. Evaluate the conditional expression (Section 5.6), pop the result off the expression stack, and mask off all but the lowest bit of the result to get the corresponding Boolean value (Section 4.1). If the Boolean value of the result is 1 (true), then execute the true part. If the Boolean value of the result is 0 (false), then either do nothing (if there is no optional ELSE clause) or execute the statement in the ELSE clause (if one is provided).

While statements. Evaluate the conditional expression (Section 5.6), pop the result off the expression stack, and mask off all but the lowest bit of the result to get the corresponding Boolean value (Section 4.1). If the Boolean value of the result is 1 (true), then execute the body of the loop and repeat the whole process. If the Boolean value of the result is 0 (false), then do nothing.

Set statements. (1) Evaluate the right-hand side expression, adjusting the size if necessary (Section 5.8) to conform to the size of the left-hand side expression. (2) The left-hand side expression must be a semantic lvalue (Section 4.7). Compute its address as specified in Section 5.7. (3) Copy the result of step (1) into memory starting at the address computed in step (2).

Call statements. Evaluate the expression contained in the statement (Section 5.6). If it produces a result of nonzero size, pop the result off the stack.

Increment statements. The operand expression must be a semantic lvalue. (Section 4.7). Compute its address as stated in Section 5.7. Add one to the value stored at the address and store the result in place.

Decrement statements. The operand expression must be a semantic lvalue. (Section 4.7). Compute its address as stated in Section 5.7. Subtract one from the value stored at the address and store the result in place.

Return statements. (1) If the statement contains an expression, then evaluate the expression (Section 5.6) and ensure that it appears at the top of the expression stack (Section 5.3) on function return. (2) Pop the current frame off the call stack and restore the parent frame.

Block statements. Execute each statement of the block in the order in which it appears in the program text.

5.6 Expressions

We define the semantics of expressions using an expression stack, as specified in Section 5.3. The compiler generates code to do the following for each program expression.

Identifiers. (1) If the identifier is a semantic lvalue (Section 4.7), then compute its address as specified in Section 5.7 and get a number of bytes equal to the size of the right-hand side (Section 4.5), starting at that address. Push those bytes on the stack, using size adjustment if necessary (Section 5.8) to generate a result that has a number of bytes equal to the left-hand expression size. (2) If the identifier corresponds to a numeric constant declaration, then push the bytes of the value associated with the declaration (Section 5.2) on the stack in little-endian order. (3) If the identifier corresponds to a data declaration, then push the two-byte address associated with the identifier (Section 5.2) on the stack in little-endian order. (4) If the identifier corresponds to a function declaration, then push the two-byte address corresponding to the first byte of the function body (Section 5.4) on the stack in little-endian order. (If the function body is defined in a different file, this can be done by using the label associated with the function body; at assembly time, the label definition will be available.)

Numeric constants. Push the bytes of the value associated with the declaration (Section 5.2) on the stack in little-endian order.

Indexed expressions. (1) Compute the address of the indexed expression as stated in Section 5.7. (2) Get a number of bytes equal to the size specified in the expression, starting at that address, and push them on the stack.

Register expressions. Push the one-byte value stored in the function call stack slot for that register (Section 5.4) on the expression stack.

Call expressions. The procedure depends on the kind of the called expression.

(1) Declared functions. If the called expression is an identifier corresponding to a function declaration, then do the following. (a) Evaluate each argument and store it to its slot on the stack in the new frame, adjusting the argument sizes (Section 5.8) as necessary to them to the parameter sizes specified in the function definition. (b) Adjust the stack pointer to point to the base of the new frame. (c) Restore all the saved registers (Section 5.4) to the values saved in their slots on the stack. (d) Do a JSR to the code for the function body (Section 5.4). (e) Save all the saved registers to their slots.

(2) *Assembly-time constants.* If the called expression evaluates to an assembly-time constant, then (a) restore the registers as in (1)(c), (b) JSR to the constant address, and (c) save the registers as in (1)(e).

(3) *Callee address unknown.* Otherwise the callee address is not known at assembly time. In this case proceed as in case 2, but in step (2)(b) store the value to zero-page memory and do a JSR to an indirect JMP to the address.

The rules are designed so that in all cases the last value written into the stack slot for each saved register prior to the call is guaranteed to be in the corresponding 6502 register just before the JSR to the called function. Thus, except for the S register, all saved registers are guaranteed to have their saved values at the point where control enters the callee function. Because of the JSR, the S register's value will be increased by 2. Similarly, all stack slots corresponding to saved registers other than the S register will contain the values that the corresponding registers had on exit from the function; the S register's value will be off by 2.

Binary operation expressions. *Shift operations* (<< and >>). Evaluate the left-hand side, evaluate the right-hand side, pop the right-hand side, and shift the left-hand side value (which is the result on the top of the stack) by a number of bytes equal to the right-hand side value. In the case of a right shift, the shift is signed (arithmetic) if the left-hand side expression is signed; otherwise it is unsigned (logical).

Arithmetic operations (*, /, +, and -). Evaluate the left and right expressions, adjust their size if necessary (Section 5.8) to make their sizes equal to the size of the entire expression, pop the results, perform the operation, and push the result. In the case of multiplication and division, the operation is signed if the entire expression is signed, otherwise unsigned. Because AppleCore uses two's complement representation for negative numbers (Section 4.1), there is no distinction between signed and unsigned addition and subtraction. The division operation produces the integer quotient. All operations are done without regard to overflow (i.e., if overflow occurs, they produce the low-order bits, up to the size of the expression).

Comparison operations (=, >, <, >=, and <=). Evaluate the left and right expressions, adjust their size if necessary (Section 5.8) to make their sizes equal to the size of the entire expression, pop the results, perform the operation, and push the result. The result is a single-byte value such that 1 represents true and 0 represents false. In the case of comparisons other than equality, the operation is signed if the entire expression is signed, otherwise unsigned.

Bitwise operations (AND, OR, and XOR). Evaluate the left and right expressions, adjust their size if necessary (Section 5.8) to make their sizes equal to the size of the entire expression, pop the results, perform the operation, and push the result.

Unary operation expressions. *Dereference* (@). The operand expression must be a semantic lvalue (Section 4.7). Compute and push its address as stated in Section 5.7.

Arithmetic negation (-) and *bitwise not* (NOT). Evaluate the operand expression and perform the operation on it. Negation is two's complement negation (i.e., NOT plus 1).

Parentheses expressions. Evaluate the expression inside the parentheses.

Typed expressions. Evaluate the value subexpression. Adjust the size to generate a new value of the correct size for the given type.

5.7 Address Computation for Semantic LValues

Computation of the address of a semantic lvalue *expr* works as follows:

1. If *expr* is an identifier corresponding to a global variable, then its address is the lowest address of the global storage assigned to that variable (Section 5.2).

2. If *expr* is an identifier corresponding to a function parameter or local variable, then its address is the lowest address of the function call stack slot assigned to that parameter or variable (Section 5.4).
3. If *expr* is an indexed expression, then its address is computed by evaluating the base expression, evaluating the offset expression, popping and adding the results, and adjusting the size of the result to 2 if necessary (Section 5.8).
4. If *expr* is a register expression, then its address is the address of the function call stack slot assigned to that register (Section 5.4).

5.8 Size Adjustment

Size adjustment applies whenever a result is assigned into or out of a semantic lvalue and the sizes of the source and target values do not match. The following rules apply to size adjustment:

1. If the source value size is greater than the target value size, then high-order bytes of the source value are *truncated*: only the low-order bytes of the source value are stored, up to the number of bytes required by the target value.
2. If the source value size is smaller than the target value size, and the source value is unsigned, then the source value is *zero extended*: the bytes of the source value form the low-order bytes of the target value, and zeros are used to fill the rest of the bytes of the target value.
3. If the source value size is smaller than the target value size, and the source value is signed, then the source value is *sign extended*: the bytes of the source value form the low-order bytes of the target value, and the rest of the bytes of the target value are filled with either all 0 bits (if the high-order bit of the source value is 0) or all 1 bits (if the high-order bit of the source value is 1).